

CS144 Notes: Spatial Index

Main question

- Many queries are location-sensitive, especially queries from mobile devices.
 - e.g., where is a nearby gas station?
- Common types of spatial queries
 - *Range queries*: Given a “range” (say, a rectangular bounding box), return the set of points within the range
 - *Nearest-neighbor queries*: Return the closest point to a given point
 - *Where-am-I queries*: Given a point, which object overlaps with the point?
 - *Partial-match queries*: Given a value for a specific dimension, return all points matching the value in that dimension
 - E.g. What countries are on equator?
- Q: How can we support spatial queries efficiently?

Spatial queries in SQL

Table Point(x, y):

- Q1: SQL query for “find all points in the rectangle $10 < x < 20$ and $30 < y < 40$ ”?
- Q2: SQL query for “find the nearest point to (10, 20)”?
- Q: How do we answer Q1? Any efficient way? B+tree?
 - Reminder
 - Data is stored in disks
 - Data is stored in the unit of “blocks” in disk

- Example:
 - 1,000,000 points within of $0 < x < 100$ and $0 < y < 100$. Uniformly distributed.
 - 100 points per disk block
 - Q: How to build indexes?

- Q: We may cluster points based on the primary (clustered) index. What about the secondary index?

- Q: How will the points be stored in disk blocks?

- Q: How many points are retrieved?

- Q: Is it helpful to have a secondary index for the given query?

- Q: How do we answer Q2? Any efficient way?
 - Q: Pick a range in each dimension. Ask the range query. Select the closest point to the target within the range. Any problem? Two possible outcome
 - Q: Case 1. No point within the selected range. What should we do?

- Q: How should we pick the initial range?

- Q: Case 2. A point found in the range. Is it the nearest point?


- Q: When do we have to worry about a closer point outside the range?

- Q: 1,000,000 points uniformly distributed between $0 < x < 1,000$ and $0 < y < 1,000$. Nearest point to (100, 200). What is a good choice for the initial range d of $100 - d < x < 100 + d$ and $200 - d < y < 200 + d$?


- Q: How can we build an index for “multidimensional” or “spatial” queries?

Grid File

- Space is partitioned into a “grid”
 - In each dimension, “gridlines” partition the space into “stripes”.
 - The number of gridlines and their spacing may vary among dimensions
- Example: (age, salary) data
 - (25, 60K), (45, 60K), (50, 75K), (50, 100K), (50, 120K), (70, 110K), (85, 140K), (30, 260K), (25, 400K), (45, 350K), (50, 275K), (60, 260K)
 - Conceptual grid file: (2 data points per bucket)



- Grid file data structure.
 - $(n \times n)$ array of bucket pointers with their range values
 - A set of data buckets storing the actual data points



- E.g., Lower left bucket stores all points in $(0 \leq x < 40, 0 \leq y < 90)$
- Querying grid file
 - Partial-match queries
 - Q: Find all customers with a salary of \$200K? What buckets need to be searched?
 - Range queries
 - Q: Find all customers aged 35-45 and with a salary of 50K-100K? What buckets need to be searched?
 - Nearest-neighbor queries
 - Q: Find the point nearest to (45, 200K)? What buckets need to be searched?
- Updating grid file
 - Q: Insert (52, 199K). What should we do when the bucket is full?

- Choice 1: Add overflow bucket. Pros and cons?


- Choice 2: Reorganize by adding or moving the grid lines.
 - Challenge: adding a grid line splits all buckets along the line
 - #1: Split with a vertical line? Where?

 - #2: Split with a horizontal line? Where?


- Generally, grid file is not suitable for a dynamic dataset with many insertions

Quad Tree

- A region is recursively partitioned into four equal-sized quadrants until all points in a region can fit into a single disk bucket.
- Conceptual quad tree example: (each bucket can store 2 points)



- Quad tree data structure:




- Each internal node (circle) splits a region into four quadrants and stores pointers to its children
 - Leaf nodes corresponds to one disk bucket and stores data points (rectangle)
- Querying a quad tree
 - Partial-match queries
 - Q: Find all customers with a salary of \$200K? What buckets need to be searched?
 - Range queries
 - Q: Find all customers aged 35-45 and with a salary of 50K-100K? What buckets need to be searched?
 - Nearest-neighbor queries
 - Q: Find the point nearest to (45, 200K)? What buckets need to be searched?
 - At each internal node, check overlap of the query region with each quadrant and search any quadrant with overlap
- Updating a quad tree
 - Q: Insert (52, 199K). What should we do when the bucket is full?

- Q: Is the tree always balanced? Is the tree “well occupied”?
- Comments: Quad tree is simple to implement yet its rigidity may lead to skews (and low performance) in the data structure
- Note:
 - A region does not need to be split into (2 x 2) quadrants. Many systems split a region into (n x n) quadrants, where n is a configuration parameter.
 - It is possible to extend the quad-tree algorithms to insert “shapes” not just “points” into the tree. We do not cover this extension in this class.
 - Oracle and Microsoft SQL Server support quad tree.


R-tree (Region tree)

- Each internal node of an R-tree contains a set of (arbitrary-sized) subregions as its children. Data points are stored at leaf nodes.
 - Example: An R-tree node region and the subregions of its children



- Note: When there is an overlap in subregions, a data point may be stored in any one of the overlapping subregions.

- An R-tree node data structure



- Querying an R-tree
 - Q: Find all customers aged 35-45 and with a salary of 50K-100K? What buckets need to be searched?

- At each internal node, check overlap of the query region with each subregion and search *all* subregions with overlap
- Updating an R-tree
 - Insert a new point.
 - Q: What if there is no subregion containing the point?
- In subregion expansion, minimal expansion is often preferred.
 - Q: What if there is no space in the corresponding subregion?
- In subregion split, we generally want the two subregions to be as small as possible while covering all data points in the region.
- Note:
 - It is possible to extend the R-tree algorithms to insert “shapes” not just “points” into the tree. We do not cover this extension in this class.
 - Oracle and MySQL support R-tree.